

# Bank Note Authentication Using Deep Neural Networks

In this notebook, I have tried to predict whether a particular instance of a bank note is fake or real.

## Dataset

---

1. This data has been taken from the UCI Machine Learning Repository
2. Data were extracted from images that were taken from genuine and forged banknote-like specimens. For digitization, an industrial camera usually used for print inspection was used
3. The final images have 400x 400 pixels. Due to the object lens and distance to the investigated object gray-scale pictures with a resolution of about 660 dpi were gained. Wavelet Transform tool were used to extract features from images
4. It consists of a total of 1373 instances

## About the features

---

1. Variance : variance of Wavelet Transformed image (continuous)
2. Skewness : skewness of Wavelet Transformed image (continuous)
3. Curtosis : curtosis of Wavelet Transformed image (continuous)
4. Entropy : entropy of image (continuous)
5. Class : class (integer - 0(Fake) or 1(Real))

The first step is to import the relevant libraries

The first step is to import the relevant libraries

```
In [15]: #Importing the relevant libraries
import numpy as np
import math
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split

%matplotlib inline
plt.rcParams['figure.figsize'] = (5.0, 4.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

%load_ext autoreload
%autoreload 2

np.random.seed(1)
```

In the following blocks of code, I have implemented two mathematical functions that we will be using as activation functions in our network

1. Sigmoid Function : It is differentiable everywhere so gradient descent approach can be used for back propagation. It also outputs probability values between 0 or 1.
2. Relu Function : There is a reduced likelihood of vanishing gradient with this function. It is also more computationally efficient than other functions as it needs to only select between (0,z) where z is the input to the function

Additionally, by methods of calculus, I have implemented two more functions

{relu\_backward, sigmoid\_backward} that calculate the differentiation value of sigmoid and relu functions respectively. These, I have used during backpropagation for weight and bias vector updation.

```
In [2]: def sigmoid(Z):  
        """  
        Implements the sigmoid activation in numpy  
  
        Arguments:  
        Z -- numpy array of any shape  
  
        Returns:  
        A -- output of sigmoid(z), same shape as Z  
        cache -- returns Z as well, useful during backpropagation  
        """  
  
        A = 1/(1+np.exp(-Z))  
        cache = Z  
  
        return A, cache
```

```
In [3]: def relu(Z):  
        """  
        Implements the RELU function.  
  
        Arguments:  
        Z -- Output of the linear layer, of any shape  
  
        Returns:  
        A -- Post-activation parameter, of the same shape as Z  
        cache -- a python dictionary containing "A" ; stored for computing the backward pass efficiently  
        """  
  
        A = np.maximum(0,Z)  
  
        assert(A.shape == Z.shape)  
  
        cache = Z  
        return A, cache
```

```
In [35]: def relu_backward(dA, cache):
        """
        Implements the backward propagation for a single RELU unit.

        Arguments:
        dA -- post-activation gradient, of any shape
        cache -- 'Z' where we store for computing backward propagation efficiently

        Returns:
        dZ -- Gradient of the cost with respect to Z
        """

        Z = cache
        dZ = np.array(dA, copy=True) # just converting dz to a correct object.

        # When z <= 0, we set dz to 0 as well.
        dZ[Z <= 0] = 0

        assert (dZ.shape == Z.shape)

        return dZ
```

```
In [36]: def sigmoid_backward(dA, cache):
        """
        Implements the backward propagation for a single SIGMOID unit.

        Arguments:
        dA -- post-activation gradient, of any shape
        cache -- 'Z' where we store for computing backward propagation efficiently

        Returns:
        dZ -- Gradient of the cost with respect to Z
        """

        Z = cache

        s = 1/(1+np.exp(-Z))
        dZ = dA * s * (1-s)

        assert (dZ.shape == Z.shape)

        return dZ
```

The neural network (with L layers) I have implemented has Relu activation functions until the (L-1)th layer and Sigmoid function at the Lth layer. We can see, from the code below, that in a for loop that goes from 1 to L-1, I have initialized each layer's W using He initialization to regularise the model so that it does not suffer from high variance eventually. Since it is alright to initialise the b vector to all zeros, I have simply done that.

```
In [6]: def initialize_parameters_deep(layer_dims):
        """
        Arguments:
        layer_dims -- python array (list) containing the dimensions of each layer in our
        network

        Returns:
        parameters -- python dictionary containing your parameters "W1", "b1", ..., "WL",
        "bL":
            WL -- weight matrix of shape (layer_dims[l], layer_dims[l-1])
            bL -- bias vector of shape (layer_dims[l], 1)

        """

        np.random.seed(1)
        parameters = {}
        L = len(layer_dims)            # number of layers in the network

        for l in range(1, L):

            #He initialization
            parameters['W' + str(l)] = np.random.randn(layer_dims[l], layer_dims[l-1])*
            math.sqrt(2./layer_dims[l-1])
            parameters['b' + str(l)] = np.zeros((layer_dims[l], 1))

            assert(parameters['W' + str(l)].shape == (layer_dims[l], layer_dims[l-1]))
            assert(parameters['b' + str(l)].shape == (layer_dims[l], 1))

        return parameters
```

In the following code, I have implemented `linear_forward` which takes as input the activation output of the previous layer ( $X = A[0]$ ), the corresponding weight and bias matrix for the current layer

We calculate  $Z = W \cdot A_{\text{prev}} + b$

```
In [7]: def linear_forward(A, W, b):
        """
        Implements the linear part of a layer's forward propagation.

        Arguments:
        A -- activations from previous layer (or input data): (size of previous layer, number of examples)
        W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)
        b -- bias vector, numpy array of shape (size of the current layer, 1)

        Returns:
        Z -- the input of the activation function, also called pre-activation parameter
        cache -- a python dictionary containing "A", "W" and "b" ; stored for computing the backward pass efficiently
        """

        Z = W.dot(A) + b

        assert(Z.shape == (W.shape[0], A.shape[1]))
        cache = (A, W, b)

        return Z, cache
```

In the following code snippet, I have implemented a function that calls the previous function (linear\_forward) based on whether the activation is sigmoid or Relu. It would call linear\_forward with activation = "relu" for a total of L-1 layers before it calls linear\_forward with activation = "sigmoid" for the Lth layer. Additionally, linear\_activation\_forward returns a cache that is a combination of linear\_cache which stores the Z matrix where  $Z = W \cdot A_{prev} + b$  and activation\_cache stores the value of A where  $A = g(Z)$  where g is our activation function (sigmoid/relu).

```
In [8]: def linear_activation_forward(A_prev, W, b, activation):
        """
        Implements the forward propagation for the LINEAR->ACTIVATION layer

        Arguments:
        A_prev -- activations from previous layer (or input data): (size of previous layer,
        number of examples)
        W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)
        b -- bias vector, numpy array of shape (size of the current layer, 1)
        activation -- the activation to be used in this layer, stored as a text string:
        "sigmoid" or "relu"

        Returns:
        A -- the output of the activation function, also called the post-activation value
        cache -- a python dictionary containing "linear_cache" and "activation_cache";
        stored for computing the backward pass efficiently
        """

        if activation == "sigmoid":
            # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
            Z, linear_cache = linear_forward(A_prev, W, b)
            A, activation_cache = sigmoid(Z)

        elif activation == "relu":
            # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
            Z, linear_cache = linear_forward(A_prev, W, b)
            A, activation_cache = relu(Z)

        assert (A.shape == (W.shape[0], A_prev.shape[1]))
        cache = (linear_cache, activation_cache)

        return A, cache
```

In the following code, I have implemented `L_model_forward`, which calls `linear_activation_forward` `L-1` times with `activation = "relu"` and one more time for the `Lth` layer with `activation = "sigmoid"`

Additionally, "cache" will store all the caches (linear cache, activation cache) for each layer. "caches" is a combination of all such "cache" which is the second parameter returned by `L_model_forward`

```

In [9]: def L_model_forward(X, parameters):
        """
        Implements forward propagation for the [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID comp
        utation

        Arguments:
        X -- data, numpy array of shape (input size, number of examples)
        parameters -- output of initialize_parameters_deep()

        Returns:
        AL -- Last post-activation value
        caches -- list of caches containing:
                    every cache of linear_relu_forward() (there are L-1 of them, indexed
                    from 0 to L-2)
                    the cache of linear_sigmoid_forward() (there is one, indexed L-1)
        """

        caches = []
        A = X
        L = len(parameters) // 2 # number of layers in the neural network
        k

        # Implementing [LINEAR -> RELU]*(L-1).
        for l in range(1, L):
            A_prev = A
            A, cache = linear_activation_forward(A_prev, parameters['W' + str(l)], paramet
            ers['b' + str(l)], activation = "relu")
            caches.append(cache)

        # Implementing LINEAR -> SIGMOID for the Lth Layer
        AL, cache = linear_activation_forward(A, parameters['W' + str(L)], parameters['b'
        + str(L)], activation = "sigmoid")
        caches.append(cache)

        assert(AL.shape == (1,X.shape[1]))

        return AL, caches

```

In the following function, `compute_cost`, we calculate the cost of a neural network model with a binary classification output of 0 or 1. From calculus, we see that :

$$\text{cost} = (1./m) * (-\text{np.dot}(Y, \text{np.log}(AL).T) - \text{np.dot}(1-Y, \text{np.log}(1-AL).T))$$

is the most efficient way to calculate the cost. If `ycap` is our prediction and `y` is our true label, we see that upon substituting `y = 1` in the above function, this cost equation forces `ycap` to be 1 in order to minimize the loss to approximately 0 and vice versa for `y = 0`.

```
In [10]: def compute_cost(AL, Y):
    """
    Implements the cost function for a binary classification problem

    Arguments:
    AL -- probability vector corresponding to your label predictions, shape (1, number of examples)
    Y -- true "label" vector (for example: containing 0 if non-cat, 1 if cat), shape (1, number of examples)

    Returns:
    cost -- cross-entropy cost
    """

    m = Y.shape[1]

    # Compute loss from aL and y.
    cost = (1./m) * (-np.dot(Y,np.log(AL).T) - np.dot(1-Y, np.log(1-AL).T))

    cost = np.squeeze(cost)      # this turns [[17]] into 17
    assert(cost.shape == ())

    return cost
```

In the following function, `linear_backward`, I have calculated `dW` and `dB` given `dZ` of that current layer and cache which stores (`A_prev`, `W` and `b`) from forward propagation for that layer.

```
In [11]: def linear_backward(dZ, cache):
    """
    Implements the linear portion of backward propagation for a single layer (layer L)

    Arguments:
    dZ -- Gradient of the cost with respect to the linear output (of current layer L)
    cache -- tuple of values (A_prev, W, b) coming from the forward propagation in the current layer

    Returns:
    dA_prev -- Gradient of the cost with respect to the activation (of the previous layer L-1), same shape as A_prev
    dW -- Gradient of the cost with respect to W (current layer L), same shape as W
    db -- Gradient of the cost with respect to b (current layer L), same shape as b
    """

    A_prev, W, b = cache
    m = A_prev.shape[1]

    dW = 1./m * np.dot(dZ,A_prev.T)
    db = 1./m * np.sum(dZ, axis = 1, keepdims = True)
    dA_prev = np.dot(W.T,dZ)

    assert (dA_prev.shape == A_prev.shape)
    assert (dW.shape == W.shape)
    assert (db.shape == b.shape)

    return dA_prev, dW, db
```



The following function `linear_activation_backward` calls the appropriate function (`relu_backward/sigmoid_backward`) based on what the activation value is (`relu/sigmoid`). Essentially, it makes this call to calculate `dA_prev` which is the differential value of the activation function of the previous layer.

```
In [12]: def linear_activation_backward(dA, cache, activation):
         """
         Implements the backward propagation for the LINEAR->ACTIVATION layer.

         Arguments:
         dA -- post-activation gradient for current layer L
         cache -- tuple of values (linear_cache, activation_cache) we store for computing
         backward propagation efficiently
         activation -- the activation to be used in this layer, stored as a text string:
         "sigmoid" or "relu"

         Returns:
         dA_prev -- Gradient of the cost with respect to the activation (of the previous L
         ayer L-1), same shape as A_prev
         dW -- Gradient of the cost with respect to W (current layer L), same shape as W
         db -- Gradient of the cost with respect to b (current layer L), same shape as b
         """
         linear_cache, activation_cache = cache

         if activation == "relu":
             dZ = relu_backward(dA, activation_cache)
             dA_prev, dW, db = linear_backward(dZ, linear_cache)

         elif activation == "sigmoid":
             dZ = sigmoid_backward(dA, activation_cache)
             dA_prev, dW, db = linear_backward(dZ, linear_cache)

         return dA_prev, dW, db
```

`L_model_backward` runs the main function for backpropagation for the whole model, where it calls `linear_activation_backward` once for the Lth layer with `activation = "sigmoid"` and otherwise L-1 times with `activation = "relu"`

```
In [13]: def L_model_backward(AL, Y, caches):
    """
    Implements the backward propagation for the [LINEAR->RELU] * (L-1) -> LINEAR -> S
    IGMOID group

    Arguments:
    AL -- probability vector, output of the forward propagation (L_model_forward())
    Y -- true "label" vector (containing 0 if note is fake, 1 if note is real)
    caches -- list of caches containing:
                every cache of linear_activation_forward() with "relu" (there are (L-
    1) or them, indexes from 0 to L-2)
                the cache of linear_activation_forward() with "sigmoid" (there is on
    e, index L-1)

    Returns:
    grads -- A dictionary with the gradients
                grads["dA" + str(L)] = ...
                grads["dW" + str(L)] = ...
                grads["db" + str(L)] = ...

    """
    grads = {}
    L = len(caches) # the number of layers
    m = AL.shape[1]
    Y = Y.reshape(AL.shape) # after this line, Y is the same shape as AL

    # Initializing the backpropagation
    dAL = - (np.divide(Y, AL) - np.divide(1 - Y, 1 - AL))

    # Lth Layer (SIGMOID -> LINEAR) gradients. Inputs: "AL, Y, caches". Outputs: "gra
    ds["dAL"], grads["dWL"], grads["dbL"]
    current_cache = caches[L-1]
    grads["dA" + str(L-1)], grads["dW" + str(L)], grads["db" + str(L)] = linear_activ
    ation_backward(dAL, current_cache, activation = "sigmoid")

    for l in reversed(range(L-1)):
        # lth Layer: (RELU -> LINEAR) gradients.
        current_cache = caches[l]
        dA_prev_temp, dW_temp, db_temp = linear_activation_backward(grads["dA" + str(
    l + 1)], current_cache, activation = "relu")
        grads["dA" + str(l)] = dA_prev_temp
        grads["dW" + str(l + 1)] = dW_temp
        grads["db" + str(l + 1)] = db_temp

    return grads
```

update\_parameters will update W and b based on the differentiation parameters that it receives after calculation and also the learning rate(alpha). It will do so for each layer in the neural network.

```
In [14]: def update_parameters(parameters, grads, learning_rate):
        """
        Updates parameters using gradient descent

        Arguments:
        parameters -- python dictionary containing your parameters
        grads -- python dictionary containing your gradients, output of L_model_backward

        Returns:
        parameters -- python dictionary containing your updated parameters
                       parameters["W" + str(L)] = ...
                       parameters["b" + str(L)] = ...
        """

        L = len(parameters) // 2 # number of layers in the neural network

        # Update rule for each parameter. Use a for loop.
        for l in range(L):
            parameters["W" + str(l+1)] = parameters["W" + str(l+1)] - learning_rate * grads["dW" + str(l+1)]
            parameters["b" + str(l+1)] = parameters["b" + str(l+1)] - learning_rate * grads["db" + str(l+1)]

        return parameters
```

load\_data\_csv : Loads data from the designated folder in "csv" format. It also splits the data as follows:  
train -> 80% test -> 20%

Eventually, we will reshape the data to the appropriate dimensions in order to pass as input to  
initialize\_parameters\_deep

```
In [25]: def load_data_csv(name_of_file):
        path = r'C:\Users\Aishwarya\Desktop\Desktop\UdemyCoursera\BankNoteAuthentication\banknote.csv'
        full_dataset = pd.read_csv(path+name_of_file)

        train, test = train_test_split(full_dataset, test_size=0.2)

        train_set_x = np.array(train.iloc[:,0:4])
        train_set_y = np.array(train.iloc[:,4])

        test_set_x = np.array(test.iloc[:,0:4])
        test_set_y = np.array(test.iloc[:,4])

        train_set_y = train_set_y.reshape((1, train_set_y.shape[0]))
        test_set_y = test_set_y.reshape((1, test_set_y.shape[0]))

        #print(train_set_x.shape)
        #print(test_set_y.shape)

        return train_set_x, train_set_y, test_set_x, test_set_y
```

```
In [26]: train_x_orig, train_y, test_x_orig, test_y = load_data_csv('banknote.csv')
```

```
In [27]: # Exploring our dataset
m_train = train_x_orig.shape[0]
m_test = test_x_orig.shape[0]

print ("Number of training examples: " + str(m_train))
print ("Number of testing examples: " + str(m_test))
print ("train_x_orig shape: " + str(train_x_orig.shape))
print ("train_y shape: " + str(train_y.shape))
print ("test_x_orig shape: " + str(test_x_orig.shape))
print ("test_y shape: " + str(test_y.shape))
```

Number of training examples: 1097  
Number of testing examples: 275  
train\_x\_orig shape: (1097, 4)  
train\_y shape: (1, 1097)  
test\_x\_orig shape: (275, 4)  
test\_y shape: (1, 275)

```
In [30]: # Reshape the training and test examples
train_x_flatten = train_x_orig.reshape(train_x_orig.shape[0], -1).T # The "-1" make
s reshape flatten the remaining dimensions
test_x_flatten = test_x_orig.reshape(test_x_orig.shape[0], -1).T

# Standardize data to have feature values between 0 and 1.
train_x = train_x_flatten/255.
test_x = test_x_flatten/255.

print ("train_x's shape: " + str(train_x.shape))
print ("test_x's shape: " + str(test_x.shape))
```

train\_x's shape: (4, 1097)  
test\_x's shape: (4, 275)

We define our network as follows:

1. Input layer : 4 neurons, each corresponding to one feature of the input observation
2. 1st Hidden layer : 20 neurons
3. 2nd Hidden layer : 100 neurons
4. 3rd Hidden layer : 75 neurons
5. 4th Hidden layer : 50 neurons
6. 5th Hidden layer : 20 neurons
7. 6th Hidden layer : 10 neurons
8. Output Layer : 1 neuron

Here, we purposely select a neural network which is very deep with so many layers, because we have very less data at hand, so in order to combat the problem of high bias, we will make the neural network more complex. Since there is very less data, we do not need to do mini-batch processing as we have enough computational memory to run around 1300 observations which would not even take much time.

```
In [19]: ### CONSTANTS ###
layers_dims = [4,20,100, 75,50,20,10,1] # 7-Layer model
```

L\_layer\_model will call the appropriate neural network functions defined above and calculate the cost of our model after every 1000 iterations. It will do so as many times as the value of num\_iterations which is the epoch value that we will input for our model.

```

In [21]: def L_layer_model(X, Y, layers_dims, learning_rate = 0.0075, num_iterations = 3000, p
rint_cost=False):
    """
    Implements a L-layer neural network: [LINEAR->RELU]*(L-1)->LINEAR->SIGMOID.

    Arguments:
    X -- data, numpy array of shape (num_px * num_px * 3, number of examples)
    Y -- true "label" vector (containing 0 if cat, 1 if non-cat), of shape (1, number
of examples)
    layers_dims -- list containing the input size and each layer size, of length (num
ber of layers + 1).
    learning_rate -- Learning rate of the gradient descent update rule
    num_iterations -- number of iterations of the optimization loop
    print_cost -- if True, it prints the cost every 100 steps

    Returns:
    parameters -- parameters learnt by the model. They can then be used to predict.
    """

    np.random.seed(1)
    costs = [] # keep track of cost

    # Parameters initialization
    parameters = initialize_parameters_deep(layers_dims)

    # Loop (gradient descent)
    for i in range(0, num_iterations):

        # Forward propagation: [LINEAR -> RELU]*(L-1) -> LINEAR -> SIGMOID.
        AL, caches = L_model_forward(X, parameters)

        # Compute cost.
        cost = compute_cost(AL, Y)

        # Backward propagation.
        grads = L_model_backward(AL, Y, caches)

        # Update parameters.
        parameters = update_parameters(parameters, grads, learning_rate)

        # Print the cost every 1000 training example
        if print_cost and i % 1000 == 0:
            print ("Cost after iteration %i: %f" %(i, cost))
        if print_cost and i % 1000 == 0:
            costs.append(cost)

    # plot the cost
    plt.plot(np.squeeze(costs))
    plt.ylabel('cost')
    plt.xlabel('iterations (per 000s)')
    plt.title("Learning rate =" + str(learning_rate))
    plt.show()


    return parameters

```

Now, we will run our model, and note the cost after every 1000 iterations. Plotting a graph of cost after every 1000 iterations shows us how our model reaches the appropriate optimal value.

```
In [31]: parameters = L_layer_model(train_x, train_y, layers_dims, num_iterations = 8000, print_cost = True)
```

```
Cost after iteration 0: 0.694189  
Cost after iteration 1000: 0.524189  
Cost after iteration 2000: 0.497194  
Cost after iteration 3000: 0.449796  
Cost after iteration 4000: 0.402558  
Cost after iteration 5000: 0.216739  
Cost after iteration 6000: 0.128143  
Cost after iteration 7000: 0.074602
```



Now, we predict our model's performance on unseen-data(test data)

```
In [32]: def predict(X, y, parameters):
        """
        This function is used to predict the results of a L-layer neural network.

        Arguments:
        X -- data set of examples you would like to label
        parameters -- parameters of the trained model

        Returns:
        p -- predictions for the given dataset X
        """

        m = X.shape[1]
        n = len(parameters) // 2 # number of layers in the neural network
        p = np.zeros((1,m))

        # Forward propagation
        probas, caches = L_model_forward(X, parameters)

        # convert probas to 0/1 predictions
        for i in range(0, probas.shape[1]):
            if probas[0,i] > 0.5:
                p[0,i] = 1
            else:
                p[0,i] = 0

        print("Accuracy: " + str(np.sum((p == y)/m)))

        return p
```

```
In [33]: pred_train = predict(train_x, train_y, parameters)
```

Accuracy: 0.9644484958979032

```
In [34]: pred_test = predict(test_x, test_y, parameters)
```

Accuracy: 0.9636363636363636

We can see that our model, performs very well on both train as well as test data. This suggests, that we have low variance and low bias and we are not overfitting our model on training data.

```
In [ ]:
```